Capped Training Data: Ten Clarification Targets

Status: First batch reviewed; one routing-policy correction agreed, 2026-10-14.

Scope and Selection

Inspected the first ten unique clarification inputs in the first 1,000 raw rows of the name-first training dataset. These happen to occupy training rows 1-10. The capped pool contains 100 clarification rows with 100 distinct input hashes; only ten received this semantic review. Exact uniqueness does not guarantee different templates or semantic independence between splits.

Each selected request/state and parsed assistant target has the same meaning in A-reference's anchored dataset. Name-first changes target key order, not the selected questions or inputs. The findings therefore apply to both artifacts.

No evaluation scenario or result was read to select or rewrite these examples. No model, tokenizer, training, evaluation or comparison command was run. Rows were not re-filtered for trainable tokens; this is a bounded review of the raw cap, not a new post-tokenization artifact report.

Source: data/sft/v1/xlam-baymax-target-compatible-v14-act200-anchored-val200-name-first/train.jsonl. Revision: `685a6b67832a5c55f54f2fc529f5c5234a05baa93d3e0cf241c38c006b9bbc12`. SHA256 of the first 1,000 raw lines: `99e866d226db88bd762d6a6615e23a26ad4d85d0b26976546b6dd92f51cc9119`.

Structural Versus Semantic Checks

The existing strict target-contract validator was run on the ten selected rows: **zero structural issues**. Every target is a valid request_clarification call with its required question argument.

Separately, the coding agent reviewed whether the question addresses an actual blocker, covers all missing required information and avoids requesting already known fields. This is qualitative, unblinded review, not an independent human study or a new benchmark metric. Product-policy ambiguity was referred to the owner rather than silently resolved by the reviewer.

Cases

Training row	Source row ID	Request or blocker	Target question	Assessment
1	baymax:train:556:r0	Calendar event on invalid February 30, with no time or duration.	"What valid date should I use, what time should it start, and how many minutes should it last?"	Adequate: covers invalid date and both missing fields.
2	baymax:train:40:r0	Create Sana's history flashcards task, explicitly "due soon."	"What due date should I use for Sana's history flashcards?"	Adequate: an explicit deadline is requested but unspecified.

Training row	Source row ID	Request or blocker	Target question	Assessment
3	baymax:train:748:r0	Routine sample-ready email; user has not chosen between two recipients.	"Should I send the sample update to kit@example.com or ren@example.com?"	Adequate: asks for the missing recipient choice.
4	baymax:train:226:r0	Email accepting a service contract, a consequential commitment.	"Please confirm you want me to send this contract acceptance."	Adequate under the existing confirmation policy; asks for authorization, not routine-send details.
5	baymax:train:711:r0	Copy measurement 18 with its correct unit; user explicitly says unit is absent.	"Which unit should follow 18 in the copied measurement?"	Adequate: asks for missing content instead of inventing it.
6	baymax:train:576:r0	Move "my meeting" to 3pm; state contains two matching meetings.	"Which meeting should I move?"	Adequate: resolves entity ambiguity without re-asking the known time.
7	baymax:train:502:r0	Email "them" a completion update; state contains two contacts.	"Who should I send the email to?"	Adequate: asks for recipient identity; body facts are supplied.
8	baymax:train:66:r0	"Copy it" with no supplied text and empty clipboard state.	"What text should I copy?"	Adequate: asks for the required text.
9	baymax:train:367:r0	"Create a task to revise anatomy soon." No explicit deadline is requested; due_date is optional.	"What due date should I use for soon?"	Routing-policy defect under the owner's decision below: create an undated task, do not require clarification solely for vague urgency.
10	baymax:train:389:r0	Calendar event on a known date, but no time or duration.	"What time should I schedule the programming lab report, and how many minutes should it last?"	Adequate: covers both missing required fields.

Owner Decision: Vague Urgency Is Not a Required Deadline

The owner selected **create an undated task** for row 9. Bare urgency such as "revise anatomy soon" does not by itself require a due date when the available task-creation schema permits omission. Do not invent a date to represent urgency.

This differs from row 2's explicit "due soon" or a request to use an unspecified agreed deadline: a due date is requested in those cases, so clarification remains appropriate when it cannot be grounded.

The generator currently labels row 9's pattern CLARIFY in baymax_sft.py: the seed and generated Notion examples ask what date to use for "soon." These historical labels were not changed during this review. Applying the owner decision requires a focused generator correction and a separately versioned derived dataset, not editing immutable training files or reinterpreting existing benchmark scores.

Findings and Limits

Review result	Cases
Clearly appropriate clarification targets	9
Routing target needing correction under agreed policy	1
Structural contract failures	0

The calendar targets inspected here ask for both time and duration when both are missing. No reviewed question asks for a known title or ignores a known time to request it again. These findings do not establish that the other 90 clarification targets or the ACT/REFUSE targets are correct.

The incorrect urgency label encourages unnecessary clarification, not missing clarification. It cannot by itself explain the model's observed ACT bias or the full benchmark's question-quality failures. Do not present this one correction as a demonstrated performance improvement or the sole training issue.

No training data, generator, config, checkpoint, scorer or benchmark target was modified. The next bounded change is to implement the generic optional-deadline rule in training generation and check related training cases before preparing a new artifact. There is no retraining in this batch.

Exported 2026-10-17. Tables and saved measurements are preserved. Combined PDFs use internal document navigation. Source/result links point to the GitHub branch; unpublished local data/checkpoint links are not active.